

Now what are the difficulties which is not yet solved and in which algorithm

The current state of transformer-based models for time-series and multimodal financial forecasting exhibits several **not fully solved difficulties**, many of which stem from both the inherent nature of time-series data and the architectural challenges of transformers. Below is a detailed summary of the main open challenges and the specific algorithms or approaches where these issues persist:

1. Quadratic Complexity and Scalability Issues

- **Problem:**
Standard transformers compute self-attention with quadratic time and memory complexity relative to the input sequence length. For long financial time-series (which can be thousands of time steps), this becomes computationally expensive and often infeasible.
- **Algorithms Affected:**
 - Vanilla Transformer (original self-attention)
 - Temporal Fusion Transformer (TFT) in large-scale settings
 - FinBERT and large LLMs for long financial texts
- **Partial Solutions:**
 - **Informer:** Uses ProbSparse attention to reduce complexity to $O(L \log L)$
 - **Performer:** Employs kernel-based approximations for linear complexity
 - **Reformer:** Uses locality-sensitive hashing (LSH) to reduce memory usageHowever, these approximations may trade off some accuracy or interpretability.

2. Data Requirements and Overfitting

- **Problem:**
Transformers generally require large datasets to learn effective representations. Financial time-series data are often noisy, non-stationary, and limited in size, leading to overfitting or poor generalization.
- **Algorithms Affected:**
 - Most transformer variants for time-series forecasting
 - FinBERT and other domain-specific LLMs when fine-tuned on small corpora
- **Challenges:**
 - Small or noisy datasets hinder transformers' ability to capture meaningful patterns

- Lack of strong inductive biases in transformers compared to traditional models (e.g., ARIMA) makes learning difficult with limited data

3. Handling Noisy, Irregular, and Multivariate Time-Series

- **Problem:**

Financial time-series are often irregularly sampled, contain missing data, and are influenced by latent variables and exogenous factors that are hard to model.

- **Algorithms Affected:**

- Vanilla transformers without special handling for irregularity
- Many time-series transformer models that assume regularly spaced inputs

- **Current Limitations:**

- Transformers struggle to incorporate irregular time intervals and missingness without explicit modeling
- Latent variable modeling remains an open problem in transformer architectures for time series

4. Lack of Strong Inductive Biases

- **Problem:**

Unlike convolutional or recurrent models, transformers have fewer built-in assumptions about temporal locality or causality, which can be a disadvantage for time-series forecasting where such biases help.

- **Algorithms Affected:**

- Standard transformers in time-series forecasting
- Pure attention-based models without hybrid convolutional or recurrent components

- **Implications:**

- Transformers may require more data and training time to learn temporal dependencies
- Hybrid models (e.g., combining CNNs or RNNs with transformers) are still an active research area

5. Multimodal Fusion Complexity

- **Problem:**

Combining heterogeneous data sources such as time-series and text remains challenging due to different sampling rates, noise levels, and feature representations.

- **Algorithms Affected:**

- Cross-modal attention fusion models
- Gated fusion networks
- Attention bottleneck fusion modules

- **Open Issues:**

- Effective temporal alignment of asynchronous textual and numerical data is non-trivial
- Fusion architectures can be computationally expensive and difficult to optimize
- Interpretability of fused representations is limited

6. Interpretability and Explainability

- **Problem:**
Transformer models, especially multimodal ones, are often black boxes, making it difficult to understand why a prediction was made.
- **Algorithms Affected:**
 - All deep transformer-based models in finance
 - Fusion models combining time-series and NLP embeddings
- **Current Limitations:**
 - Attention weights are not always reliable indicators of feature importance
 - Explainability tools like SHAP and LIME are computationally intensive on large transformers and may not fully capture complex interactions

7. Real-Time Inference and Latency

- **Problem:**
Deploying transformer models for real-time stock prediction requires ultra-low latency inference, which is challenging for large models.
- **Algorithms Affected:**
 - Large LLMs (FinBERT, LLaMA variants)
 - Deep time-series transformers with multiple layers and heads
- **Challenges:**
 - Balancing model complexity and inference speed
 - Handling streaming data synchronization and prediction under strict time constraints

Summary Table of Difficulties vs. Algorithms

Difficulty	Description	Affected Algorithms / Models	Partial Solutions / Remarks
Quadratic Complexity	High compute/memory for long sequences	Vanilla Transformer, TFT, FinBERT	Informer, Performer, Reformer reduce cost
Data Requirements & Overfitting	Need large, clean datasets	Most time-series transformers, FinBERT	Data augmentation, transfer learning
Irregular & Noisy Data Handling	Missing data, irregular sampling	Vanilla transformers	Specialized preprocessing, hybrid models
Lack of Inductive Biases	Weak temporal assumptions	Pure transformer models	Hybrid CNN/Transformer or RNN/Transformer

Difficulty	Description	Affected Algorithms / Models	Partial Solutions / Remarks
Multimodal Fusion Complexity	Aligning and fusing heterogeneous modalities	Cross-modal attention, gated fusion	Research ongoing on efficient fusion methods
Interpretability	Black-box nature, unreliable attention explanations	All deep transformer models	SHAP, LIME, attention visualization
Real-Time Latency	High inference latency in production	Large LLMs, deep transformer stacks	Model pruning, quantization, ONNX, TorchScript

References in Context

- GeeksforGeeks ^[1] highlights the quadratic complexity and model design complexity as major issues in time-series transformers.
- Reddit discussion ^[2] stresses the challenge of noisy, irregular, and limited data in time-series forecasting and the need for stronger inductive biases.
- Scientific reviews ^{[3] [4] [5]} emphasize that while transformers have revolutionized forecasting, handling irregularities, interpretability, and efficient fusion remain open research areas.
- New architectures like AutoFormer-TS ^[6] attempt to automate architecture search but do not fully solve data scarcity or interpretability.

Conclusion

While transformer architectures have advanced financial time-series and NLP modeling significantly, **key challenges remain unsolved** in:

- Efficiently scaling to very long sequences without loss of accuracy
- Robustly handling noisy, irregular, and limited data typical of financial markets
- Designing fusion architectures that effectively and interpretably combine multimodal inputs
- Ensuring real-time inference feasibility for live trading systems
- Providing trustworthy model explanations for regulatory and risk management purposes

Addressing these challenges requires ongoing research into novel architectures, better data augmentation, hybrid models, and explainability techniques tailored to financial multimodal data.

If you want, I can also provide pointers to the latest papers tackling these challenges or suggest experimental setups to explore these open problems.

✱

1. <https://www.geeksforgeeks.org/deep-learning/transformer-for-time-series-forecasting/>

2. https://www.reddit.com/r/MachineLearning/comments/1k63r4a/d_is_my_take_on_transformers_in_time_series/

3. <https://www.sciencedirect.com/science/article/pii/S2665963824001040>
4. https://isip.piconepress.com/courses/temple/ece_8110/lectures/2024_00_spring/lecture_36a.pdf
5. <https://www.mdpi.com/1424-8220/25/3/652>
6. <https://arxiv.org/abs/2502.13721>